

# Data Representation & External Sorting

DSCI 551

Wensheng Wu

# Outline

- Representing data



- How are tables stored on storage devices?

- External Sorting

- How to sort 1TB data using 1GB of memory?


In-place sorting

# Representing Data Elements

## Base address (B)


- Relational database elements:  
CREATE TABLE Product (  
pid INT PRIMARY KEY, // tinyint, smallint, mediumint  
name CHAR(20),  
description VARCHAR(100),  
? —→ maker CHAR(10) REFERENCES Company(name))
- A tuple/row is stored as a "record"

# Record Formats: Fixed Length



- Information about field types is the same for all records in a file; stored in *system catalogs*.
- **Note the importance of schema information!**

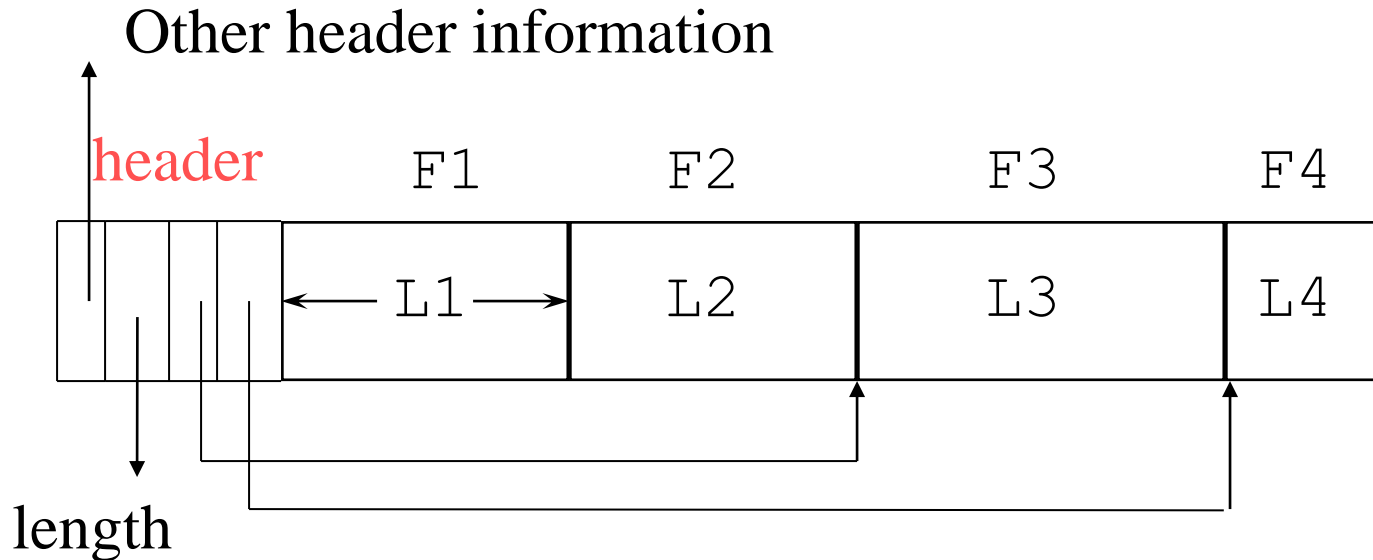
# Record Header



Header:

- Pointer to schema: help finding fields
- Length: so we know where the record ends w/o consulting schema
- Timestamp: time when record last modified or read

# Variable Length Records



Place the fixed fields first: F1, F2

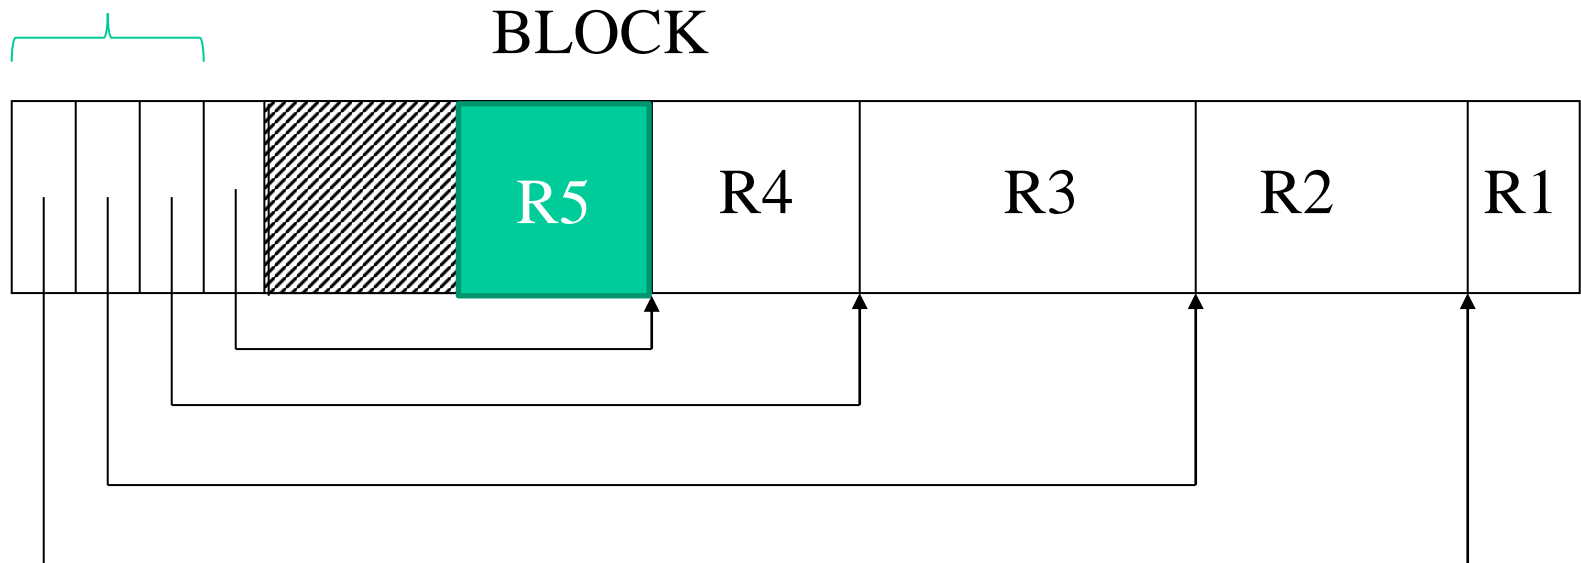
Then the variable length fields: F3, F4

Note: actually no need for pointer to F3, why?

# Storing Records in Blocks

- Blocks have fixed size (typically 4KB)
  - But records may have variable-length

Offset table (slot directory)




Why are records placed from the end?

# Problem with this design?

- Records right after slot directory
- Free space after all records

Offset table (slot directory)





# Outline

Mergesort(n)

- Representing data

- How are tables stored on storage devices?

- External Sorting



- How to sort 1TB data using 1GB of memory?

- 1GB => memory => 1GB run (sorted)

- ...

- 1GB => memory => 1GB run

} Sorting (1024)


- merging: load one block from each run to memory

# Notes

1GB



memory

---



# Notes

Memory  
1GB  
4MB  
(input buffer)



# The I/O Model of Computation

- In main memory algorithms:
  - we care about CPU time
- In databases
  - time is dominated by I/O cost
- Assumption: cost is given only by I/O
- Consequence: need to redesign certain algorithms, e.g., sorting

# Notes


- A block on storage devices loaded into a **page** in main memory
  - We sometimes interchange page with block
- Buffer pages
  - Often refer to pages in main memory used to store input, output, and intermediate data for an algorithm
- Run: a **sorted sublist** of input data

# Notes

- Make a pass through data:
  - Loading the entire data from disk once


```
select bar, count(*)  
from Sells  
group by bar
```

# Sorting

- Illustrates the difference in algorithm design when your data is not in main memory:
    - Problem: sort 1TB of data with 1GB of RAM
- R: 20 20 20 22 22 25 25 25  
S: 20 20 23 24 25 26
- Arises in many places in database systems:
    - Data requested in sorted order (ORDER BY)
    - Needed for grouping operations // group by age
    - First step in sort-merge join algorithm (R join\_a S)
    - Duplicate removal
    - Bulk loading technique for creating B+-tree indexes
- 

# 2-Way Merge-sort: Requires 3 Buffers

- Pass 0: Read a page, sort it, write it
  - only one buffer page is used
- Pass 1, 2, ..., etc.: merging two runs at a time
  - three buffer pages used.




M=3



# 2-Way Merge-sort: Requires 3 Buffers

- Pass 0: Read a page, sort it, write it
  - only one buffer page is used
- Pass 1, 2, ..., etc.: merging two runs at a time
  - three buffer pages used.



# Two-Way External Merge Sort

- Each pass we read + write each page in file.

- N pages in the file => the number of passes

$$= \lceil \log_2 N \rceil + 1$$

- So total cost is:

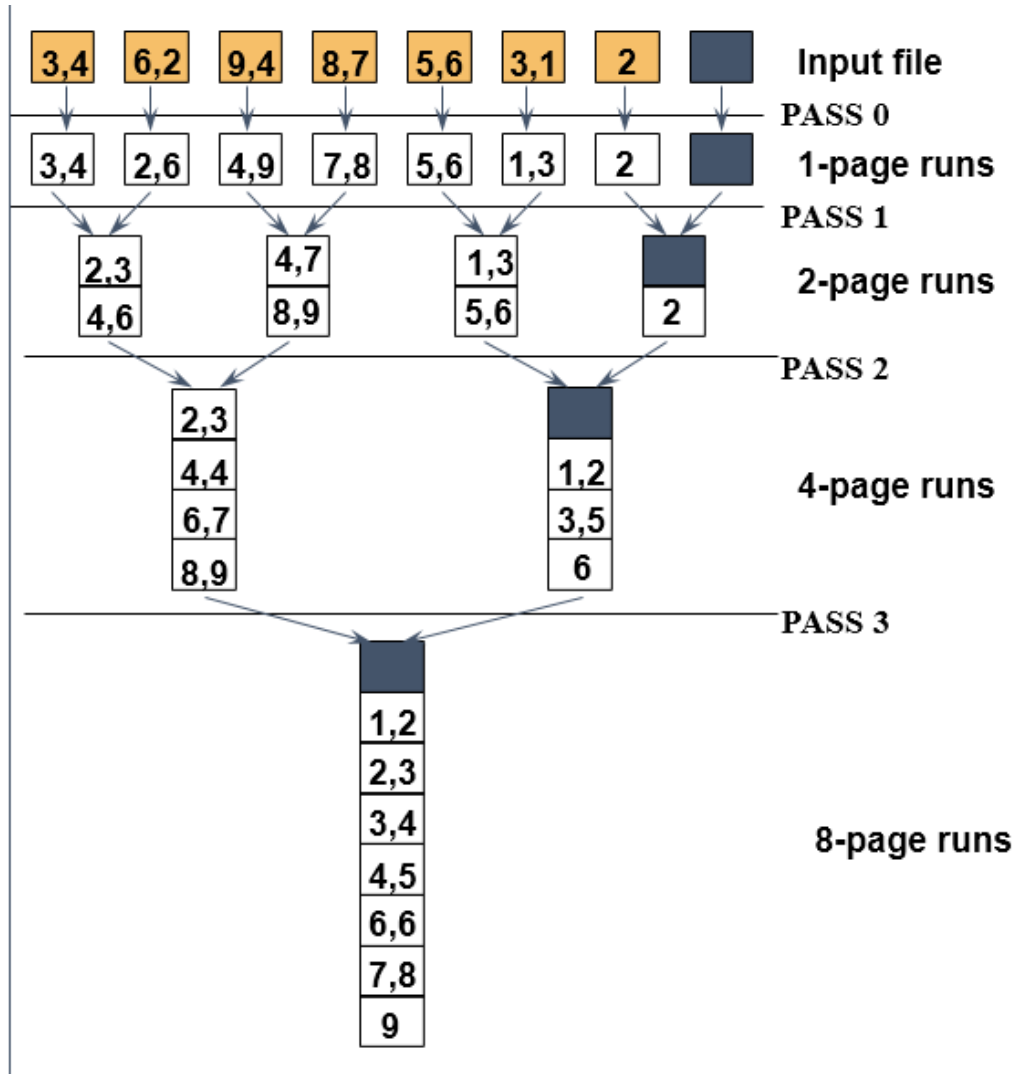
$$2N(\lceil \log_2 N \rceil + 1)$$

- Sort 4MB with buffer page size = 4KB: needs 11 passes

$$N=7, k=3$$

$$1 * 2 * 2 * \dots * 2 = 1 * 2^k \geq N \Rightarrow k \geq \log_2(N)$$

$$k = \text{ceil}[\log_2(N)] = \text{ceil}[2.8] = 3$$



# Notes

$$\begin{aligned} N &= 4\text{MB}/4\text{KB} \\ &= 1024 \\ &= 2^{10} \end{aligned}$$

$$\log_2(2^{10}) = 10$$

**M=5**,  $B(R)=108$

sorting:

# runs:  $108/5 = 21.6 = 22$

size of run:

5 blocks for each of first 21 runs

3 blocks in last run

merging 1: (4-way)

1. take 4 runs, merge into a single run

size of run:  $4*5 = 20$  blocks

2. take next 4 runs  $\Rightarrow$  20 blocks

3. next 4  $\Rightarrow$  20 blocks

4. next 4  $\Rightarrow$  20 blocks

5. next 4  $\Rightarrow$  20 blocks

6. take last two runs  $\Rightarrow 5+3 = 8$  blocks

**M=5**, B(R)=108

sorting:

# runs:  $108/5 = 21.6 = 22$

size of run:

5 blocks for each of first 21 runs

3 blocks in last run

merging 1: (4-way)

# of runs:  $6 = \text{ceiling}(22/4) = \text{ceiling}(5.5)$

size of run:

20 blocks: first 5

8 blocks: last run

merging 2: (4-way)

# of runs:  $2 = \text{ceiling}(6/4) = 2$

1. take first 4 runs  $\Rightarrow 20 * 4 = 80$  blocks

2. take last 2 runs  $\Rightarrow 20 + 8 = 28$  blocks

merging 3:

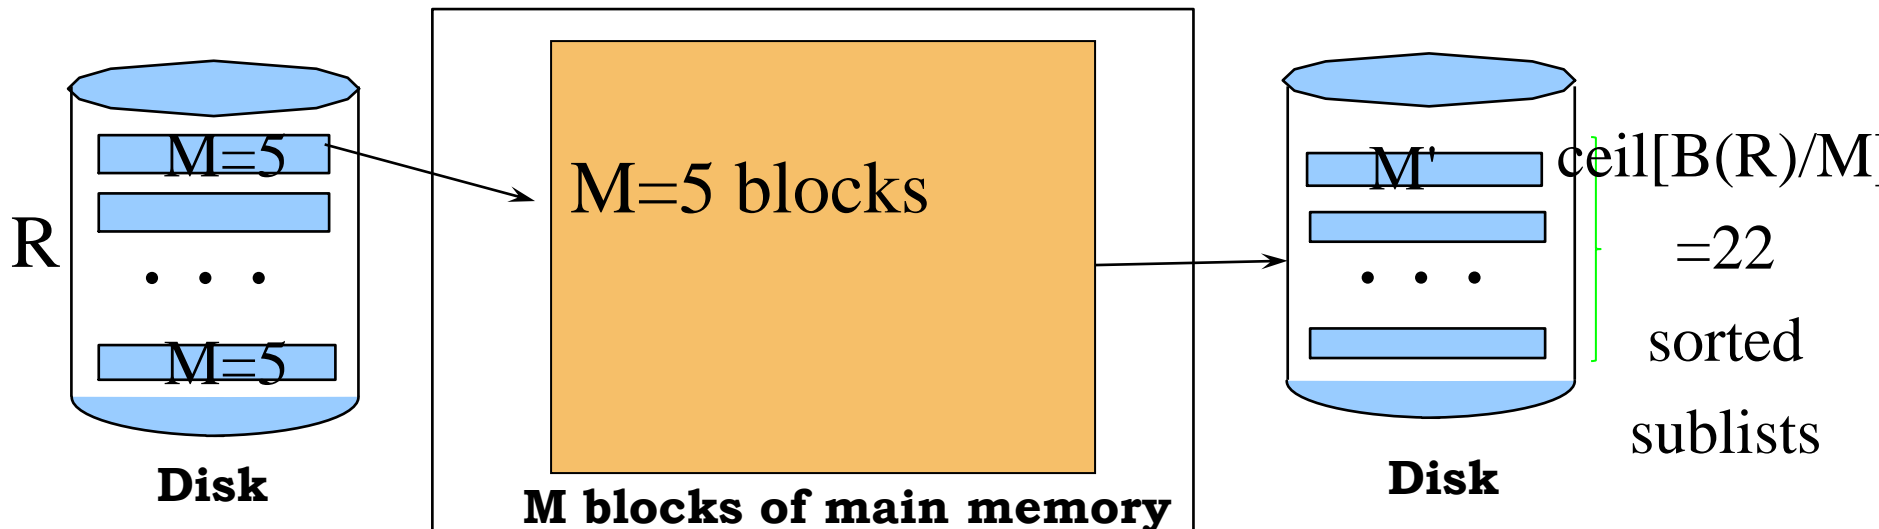
take 2 runs  $\Rightarrow$  a single run

# Can We Do Better ?

- We have more main memory
- Should use it to improve performance
- M: # of blocks (i.e., pages) in main memory
- B(R): # of blocks of relation R
  - $M=5$ ,  $B(R)=108$
  - sorting: load 5 pages, sort them, write back as run
    - $108/5 = \text{ceil}(21.6)$  runs
    - 21 runs, 5 pages/run  $\Rightarrow 21*5 = 105$  pages
    - 1 run, 3 pages/run
  - Merging: (M-1)-way merging  $\Rightarrow$  4-way merging
    - take 4 runs, 5 pages/run  $\Rightarrow$  20-page run
    - how many runs?

# External Merge-Sort

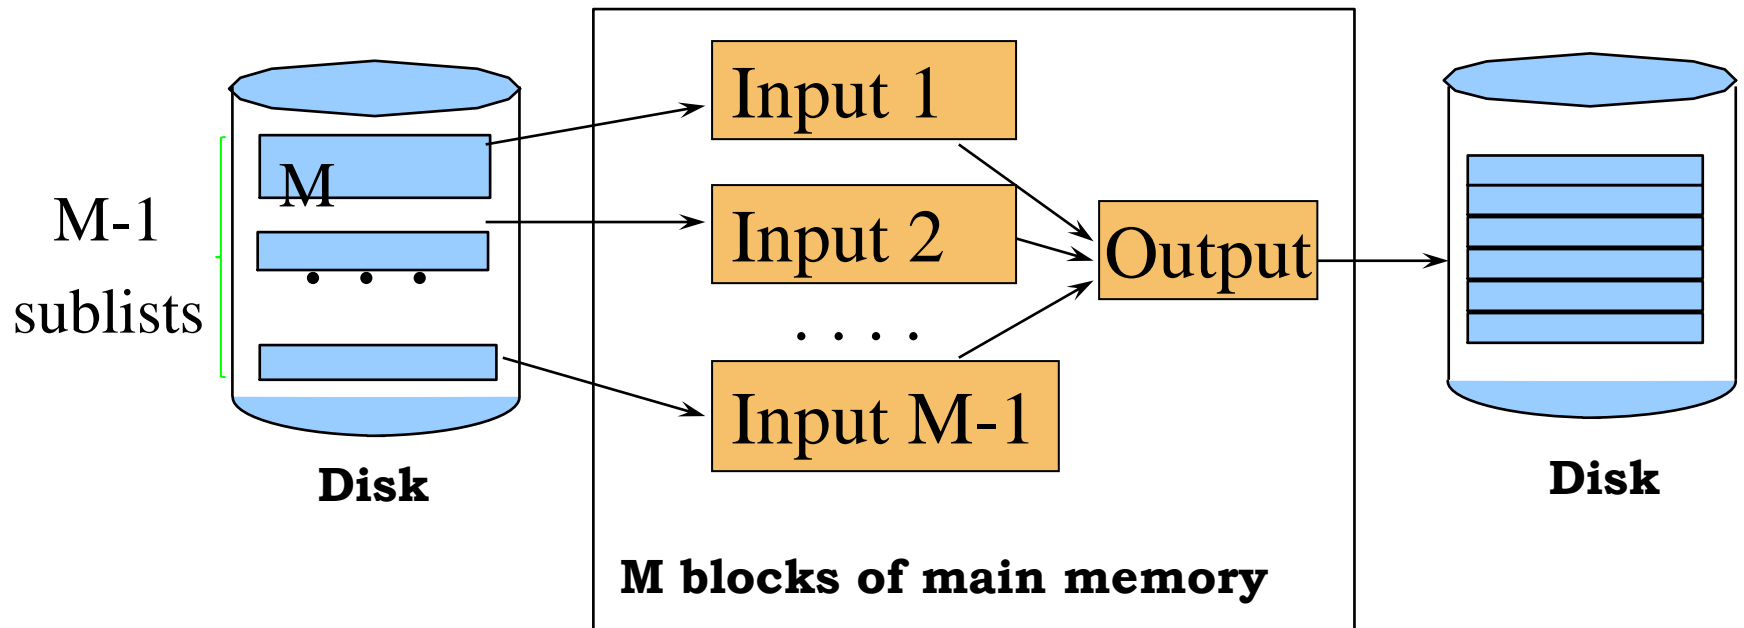
- Pass 0: load  $M$  blocks in memory, sort
  - Result:  $\text{ceil}(B(R)/M)$  sorted sublists of size  $M$
  - Each sorted sublist is a run



$$B(R)=110$$

# Pass One (merging)

- Merge  $M - 1$  runs into a new run
- Result: each run has now  $M (M - 1)$  blocks



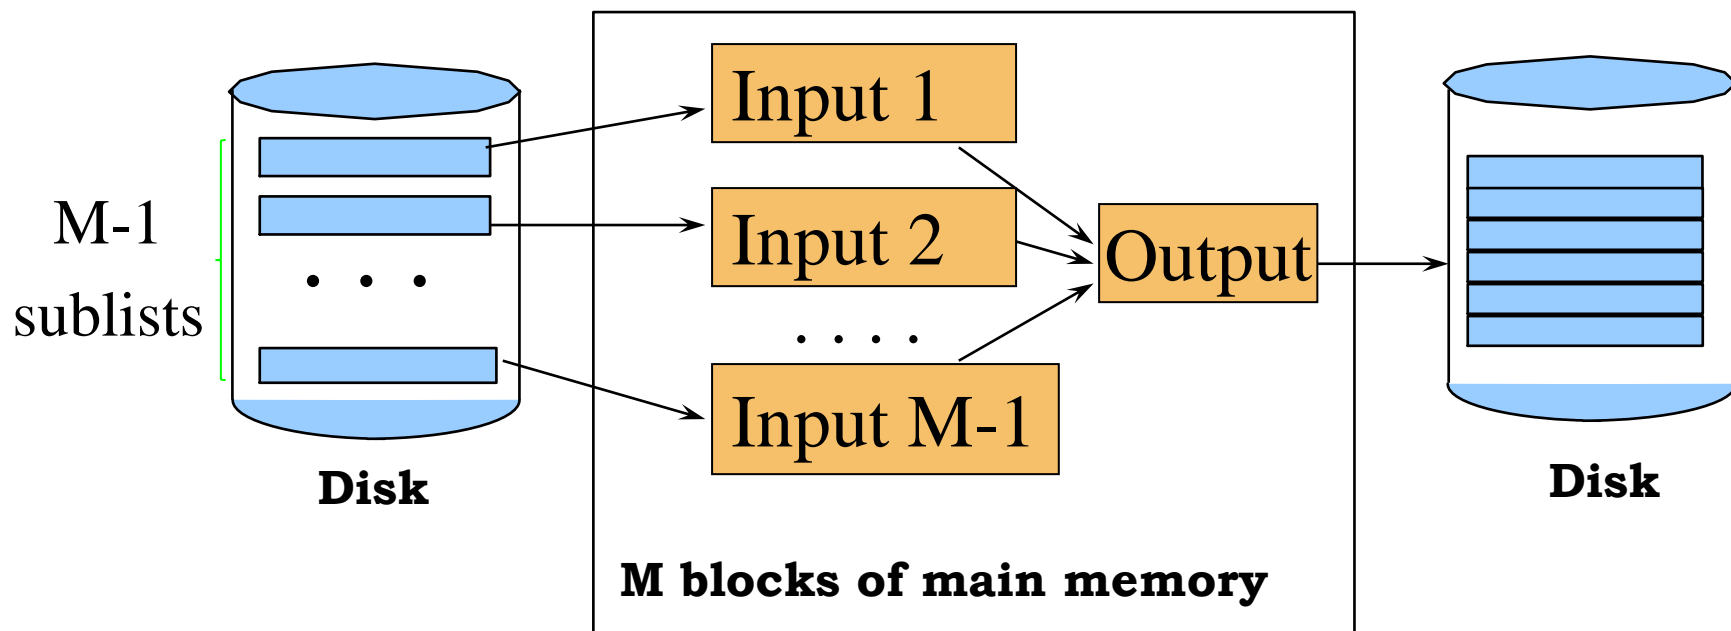


# Cost of Two-Pass, Multiway Merge Sort

- Pass 0: sort  $B/M$  sublists of size  $M$ , write
  - Cost:  $2B(R)$
- Pass 1: merge  $B/M$  sublists, write
  - Cost:  $2B(R)$
- Total cost:  $4B(R)$
- Assumption:  $B(R) \leq M^2$ 
  - $B/M \leq M - 1$  or
  - $B \leq M(M-1) \sim M^2$

# Generalized to k Passes


- Merge every  $M - 1$  runs into a new run
- Result: each run has now  $M (M - 1)^k$  blocks

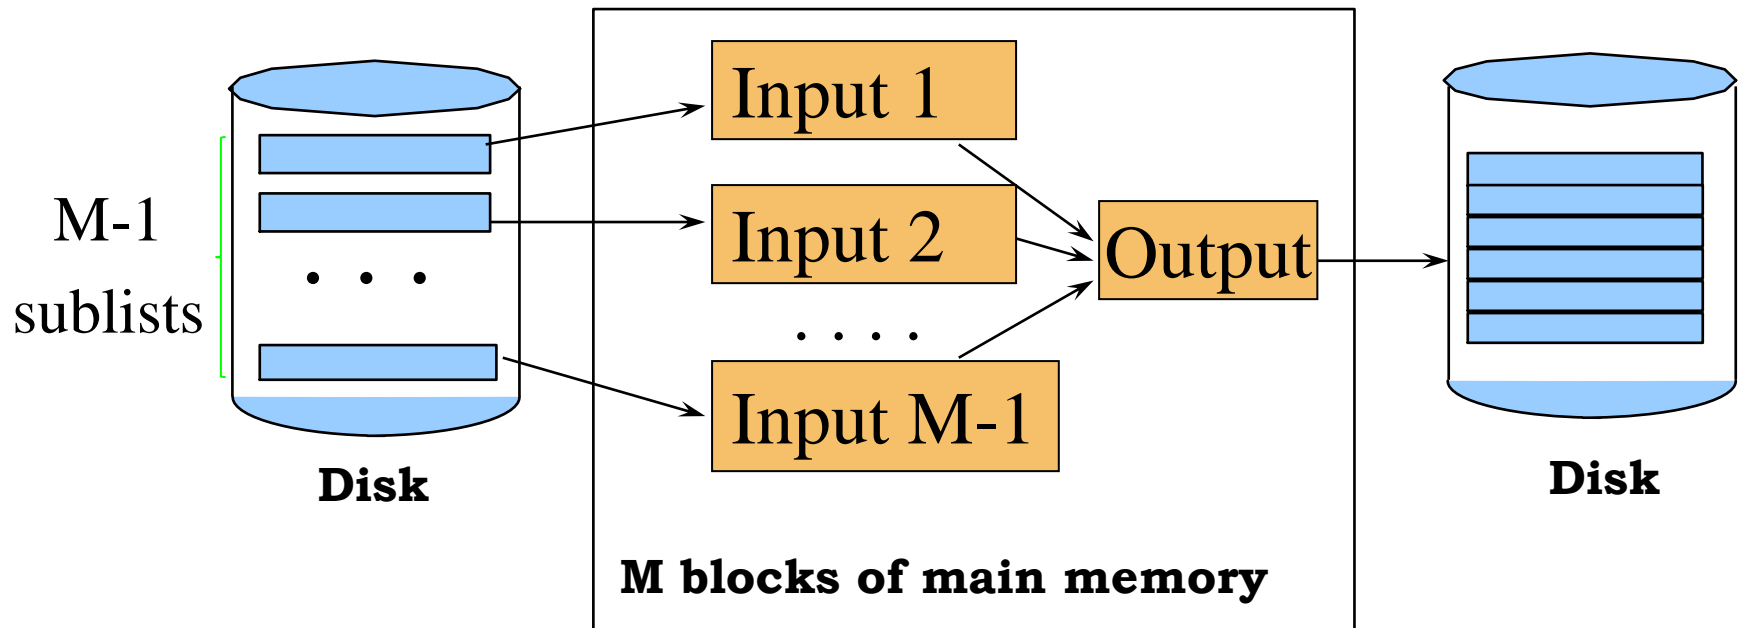


# If $k$ is the last pass

$$1 * 2^k \geq N$$

- Merge  $M - 1$  runs into a single run

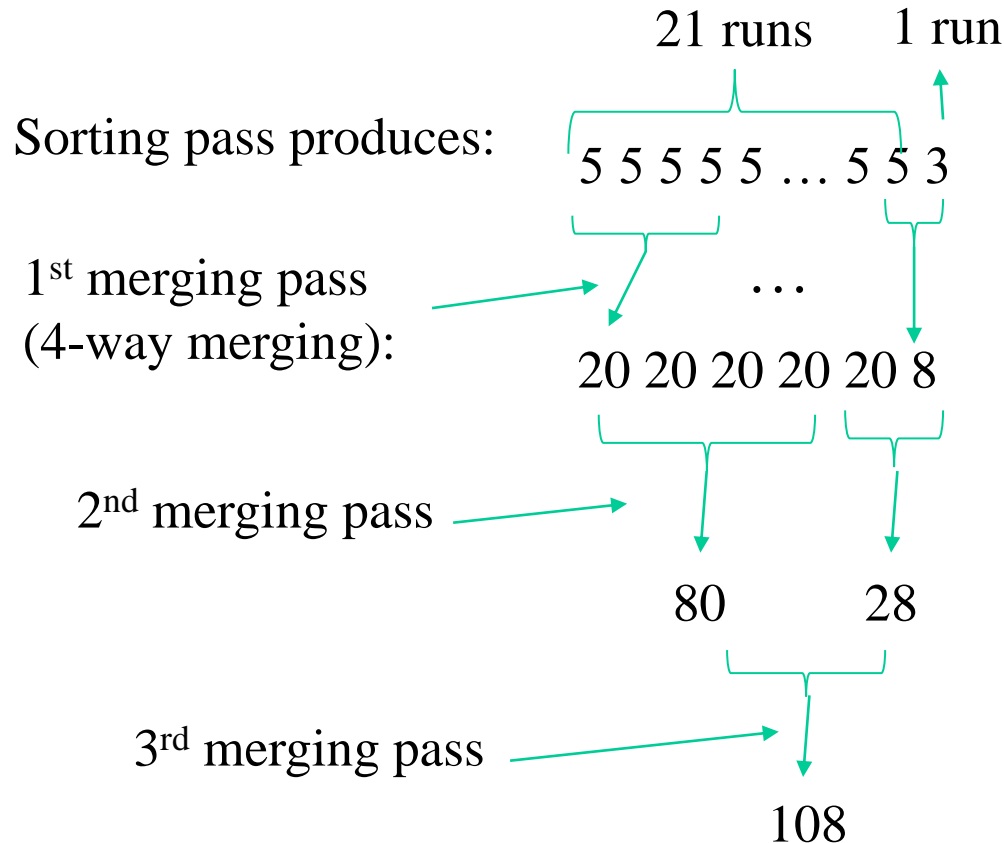
- We must have  $M (M - 1)^k \geq B(R)$    $k = \lceil \log_{M-1} \lceil B / M \rceil \rceil$



# Cost of External Merge Sort

- Number of passes:  $1 + \lceil \log_{M-1} \lceil B / M \rceil \rceil$
- Cost =  $2B * (\# \text{ of passes})$        $M=5, B(R)=108$
- E.g., with 5 buffer pages, to sort 108-page file:
  - Pass 0: produces  $\lceil 108/5 \rceil = 22$  runs (21 sorted runs of 5 pages each + last run of only 3 pages)
  - Pass 1:  $\lceil 22/4 \rceil = 6$  (5 sorted runs of 20 pages each + last run of only 8 pages)
  - Pass 2: 2 sorted runs, 80 pages and 28 pages
  - Pass 3: Sorted file of 108 pages


# Example Illustrated



# of passes:  
 $1 + \text{ceil}(\log_4(N))$

$N = \# \text{ of runs by sorting}$   
 $= 22$   
 $= \text{ceil}(108/5)$   
 $= \text{ceil}(B(R)/M)$

# Example



# Sorting 1TB using 1GB Memory

- $B(R) = 1\text{TB}/4\text{KB} = 256\text{M}$  (blocks),  $M = 1\text{GB}/4\text{KB} = 256\text{K}$  (pages)
- Sorting phase produces 1024 runs = 1K runs
- Merging:
  - Can do:  $1\text{GB}/4\text{KB}-1 = 256\text{K}-1$  ways of merging
  - Can we finish merging in one merging pass?

# # of passes

- 2-way:  $1 + \log_2(N)$ 
  - $m = 3$
  - $m' = 1$
  - $B = N$
- $(m-1)$ -way of merging:
  - $1 + \log_{(m-1)} (B/m')$