

Views

DSCI 551

Wensheng Wu

Example (ETL)

- insert into W(time, tweet, stock)
- select T.time, tweet_content, stock * 200
- from ds1.T join ds2.S on T.time = S.timestamp

- Create materialized view W(time, tweet, stock)
- As select T.time, tweet_content, stock * 200
- from ds1.T join ds2.S on T.time = S.timestamp
- <refreshing policy>

Views

- A view is a “virtual table,” a relation that is defined in terms of the contents of other tables and views.
- Declare by:

```
CREATE VIEW <name> AS <query>;
```
- In contrast, a relation whose value is really stored in the database is called a *base table*.

Example: View Definition

- CanDrink(drinker, beer) is a view “containing” the drinker-beer pairs such that the drinker frequents at least one bar that serves the beer.
- Recall Frequents(drinker, bar), Sells(bar, beer, price)

```
CREATE VIEW CanDrink AS
SELECT distinct drinker, beer
FROM Frequents, Sells
WHERE Frequents.bar = Sells.bar;
```

Example: Accessing a View

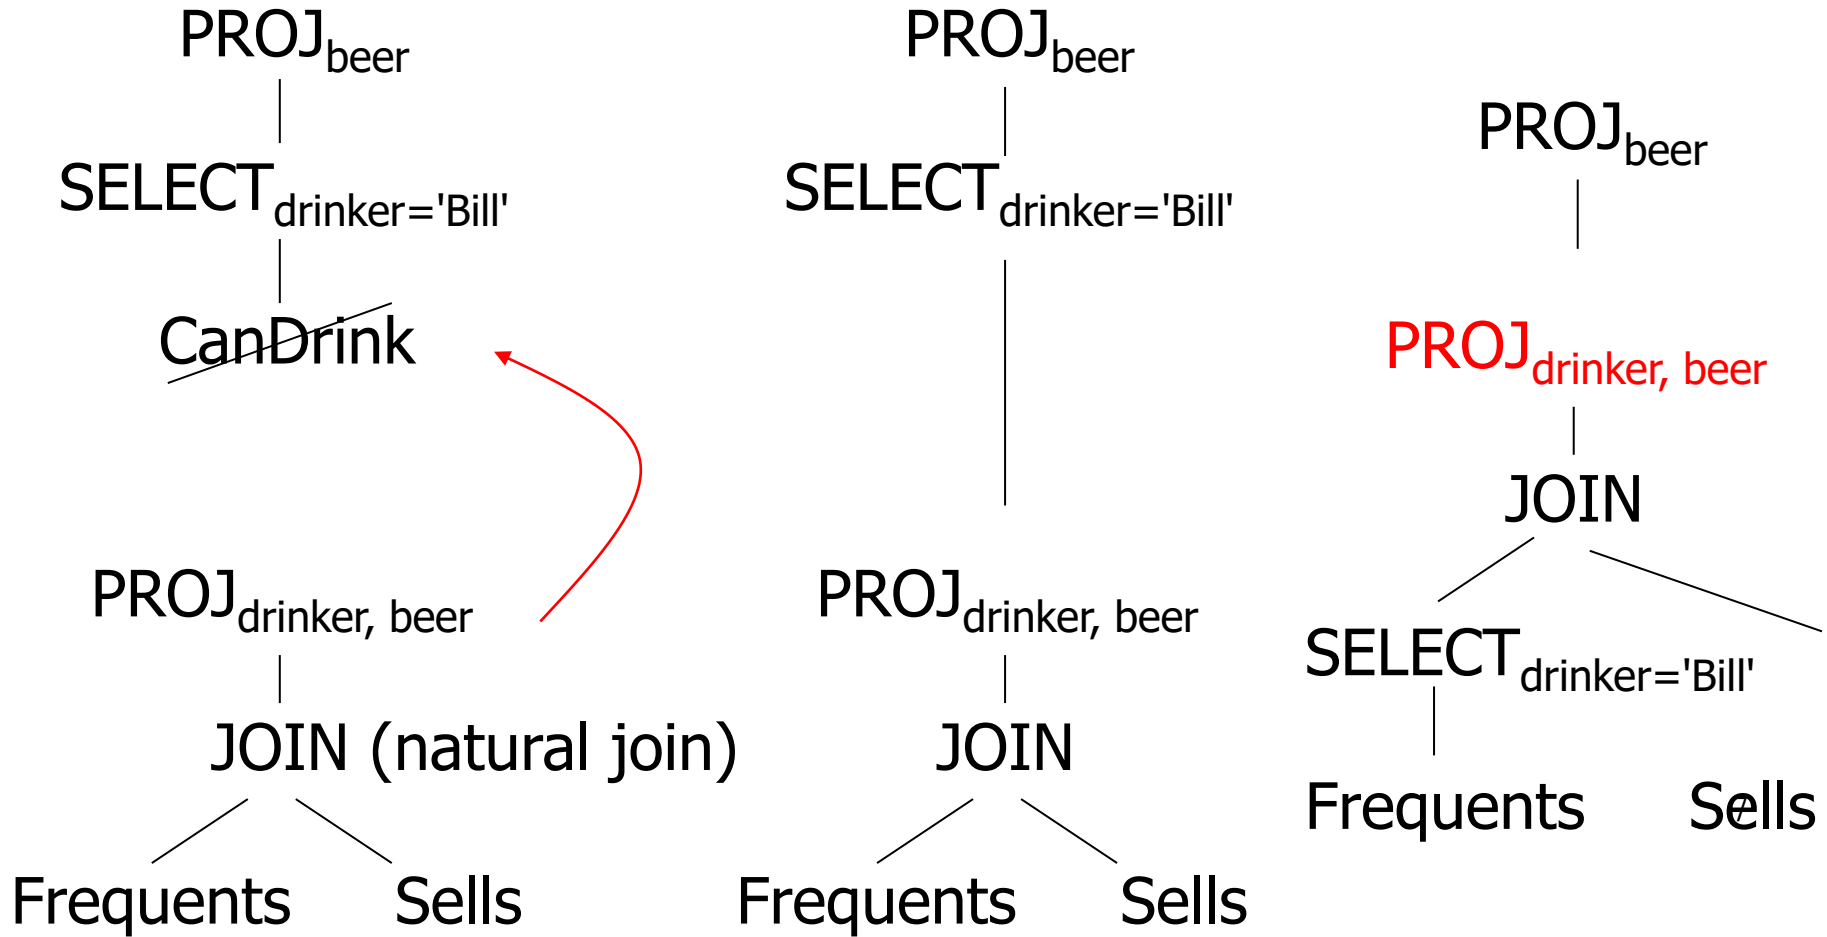
- You may query a view as if it were a base table.
 - There is a limited ability to modify views if the modification makes sense as a modification of the underlying base table.
- Example:

```
select beer from CanDrink  
where drinker = 'Bill';
```

What Happens When a View Is Used?

- The DBMS starts by interpreting the query as if the view were a base table.
 - Typical DBMS turns the query into something like relational algebra.
- The queries defining any views used by the query are also replaced by their algebraic equivalents, and “spliced into” the expression tree for the query.

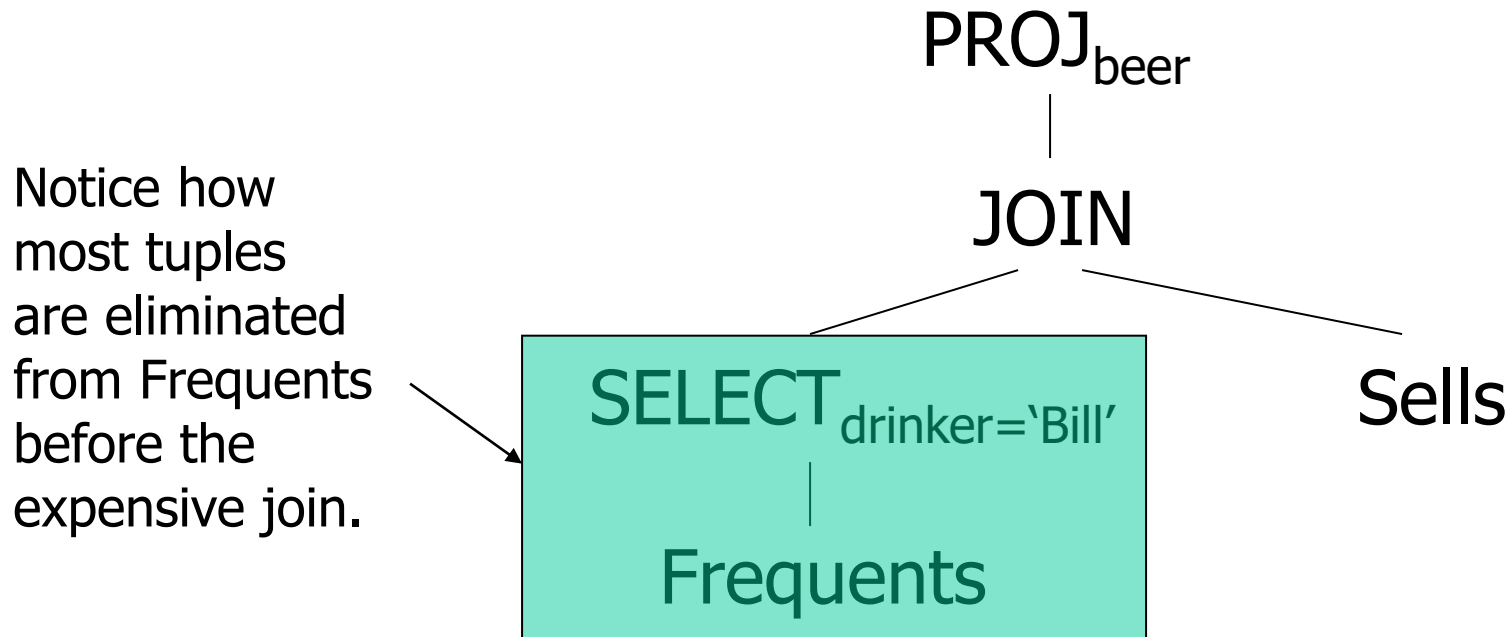
Example: View Expansion



DMBS Optimization

- It is interesting to observe that the typical DBMS will then “optimize” the query by transforming the algebraic expression to one that can be executed faster.
- Key optimizations:
 1. Push selections down the tree.
 2. Eliminate unnecessary projections.

Example: Optimization



More Examples: Defining Views

Views are relations, except that they are not physically stored.

Can be used for presenting different information to different users

Employee(ssn, name, department, project, salary)

```
CREATE VIEW Developers AS  
  SELECT name, project  
  FROM Employee  
  WHERE department = 'Development'
```

Payroll has access to all **Employees**, others only to **Developers**

A Different View

Purchase(buyer, seller, product, store, price) 1000 rows

Product(name, maker, category) 100 rows

Person(name, city, phone)

|Purchase **Join** Product| = 900?

```
CREATE VIEW LA-view AS
```

```
SELECT buyer, seller, product, store
```

```
FROM Person, Purchase
```

```
WHERE Person.city = 'LA' AND  
       Person.name = Purchase.buyer
```

We have a new virtual table:

LA-view(buyer, seller, product, store)

A Different View

LA-view(buyer, seller, product, store)

We can later use the view:

```
SELECT name, store
FROM LA-view, Product
WHERE LA-view.product = Product.name AND
      Product.category = 'shoes'
```

What Happens When We Query a View ?

Recall: LA-view(buyer, seller, product, store)

```
SELECT name, LA-view.store
FROM   LA-view, Product
WHERE  LA-view.product = Product.name AND
       Product.category = 'shoes'
```



View expansion

```
SELECT name, Purchase.store
FROM   Person, Purchase, Product
WHERE  Person.city = 'LA' AND
       Person.name = Purchase.buyer AND
       Purchase.product = Product.name AND
       Product.category = 'shoes'
```

Types of Views

- Virtual views:
 - Computed only on-demand – slow at runtime
 - Always up to date
- Materialized views
 - Precomputed offline – fast at runtime
 - Common in data warehouses (data cube)
 - Fact table + dimension tables
 - May have stale data

Reusing a Materialized View

- Suppose I have **only** the result of LAView:

```
SELECT buyer, seller, product, store
FROM   Person, Purchase
WHERE  Person.city = 'LA'   AND
       Person.name = Purchase.buyer
```

- and I want to answer the query

```
SELECT buyer, seller
FROM   Person, Purchase
WHERE  Person.city = 'LA'   AND
       Person.name = Purchase.buyer AND
       Purchase.product='gizmo'
```

Can I answer the query using only the view?

Query Rewriting Using Views

Rewritten query:

```
SELECT buyer, seller
FROM    LAView
WHERE   product= 'gizmo'
```

Original query:

```
SELECT buyer, seller
FROM    Person, Purchase
WHERE   Person.city = 'LA'   AND
        Person.name = Purchase.buyer AND
        Purchase.product='gizmo'.
```



Another Example

- I still have **only** the result of LAView:

```
SELECT buyer, seller, product, store
FROM   Person, Purchase
WHERE  Person.city = 'LA'   AND
       Person.name = Purchase.buyer
```

- but I want to answer the query

```
SELECT buyer, seller
FROM   Person, Purchase
WHERE  Person.city = 'LA'   AND
       Person.name = Purchase.buyer AND
       Person.phone LIKE '206 543 %'
```

And Now?

- I still have **only** the result of (slightly different) LAView:

```
SELECT buyer, seller, product, store
FROM Person, Purchase, Product
WHERE Person.city = 'LA' AND
      Person.name = Purchase.buyer AND
      Purchase.product = Product.name
```

- but I want to answer the query

```
SELECT buyer, seller
FROM Person, Purchase
WHERE Person.city = 'LA' AND
      Person.name = Purchase.buyer
```

And Now?

- I still have **only** the result of view SBS:

```
SELECT seller, buyer, Sum(Price) sp
FROM Purchase
WHERE Purchase.store = 'The Bon'
Group By seller, buyer
```

```
Select seller, sum(sp)
From SBS
Group by seller
```

- but I want to answer the query

```
SELECT seller, Sum(Price)
FROM Purchase
WHERE Purchase.store = 'The Bon'
Group By seller
```

And what if it's the other way around?

Example (OLAP)

Materialized view SBS(seller, buyer, sp)

| Seller | Buyer | Sum(price) sp |
|--------|----------|---------------|
| David | Bill | 10 |
| David | Jennifer | 20 |
| David | Steve | 10 |
| Bill | David | 20 |
| Bill | Mary | 10 |

Roll-up

Query:

| Seller | Sum(price) |
|--------|------------|
| David | ? |
| Bill | ? |

Finally...

- I still have **only** the result of:

```
SELECT seller, buyer, Count(*) cnt
FROM Purchase
WHERE Purchase.store = 'The Bon'
Group By seller, buyer
```

- but I want to answer the query

```
SELECT seller, Count(*)
FROM Purchase
WHERE Purchase.store = 'The Bon'
Group By seller
```

Example

View SBC(seller, buyer, cnt)

| Seller | Buyer | count(*) cnt |
|--------|----------|--------------|
| David | Bill | 2 |
| David | Jennifer | 4 |
| David | Steve | 2 |
| Bill | David | 5 |
| Bill | Mary | 2 |

Query:

| Seller | count(*) |
|--------|----------|
| David | ?? |
| Bill | ? |

Select se
From SB
Group by